SFWR ENG 4AA4

Kemal Ahmed
Fall 2015
Dr. Down

Note: information from the pre-requisite, SFWR ENG 3DX4 will not be included in this summary (although corrections will be).

Contents

Real-Time Systems

Classifications

What happens upon failure to meet deadlines:

- **Soft**: performance is degraded but not destroyed
- **Firm**: a few times will simply degrade performance, but after may lead to system failure
- **Hard**: complete and catastrophic system failure
  - **Safety Critical**: may cause injury / death (a type of hard)

Forward difference method: derivatives using $f'(x) = \frac{f(x+h) - f(x)}{h}$

Backwards Difference method: derivatives using $f'(x) = \frac{f(x) - f(x-h)}{h}$

Controller [C(s)]:
Input [E(s)]:
Output [U(s)]:
$U(s) = C(s)E(s)$

Task optimization

Task [T]: $T_i = (p_i, r_i, e_i, d_i)$
Period [p]: time between tasks are repeatedly released
Release time [r]: time it takes to release task
Execution time [e]: slowest time task could take to be completed (but assume the tasks will take this long no matter what)
Deadline [d]: when task <u>needs</u> to be completed

Number of tasks [n]:

Processor Utilization [U]: used as a priority level $U = \sum_{i=1}^{n} \frac{e_i}{p_i}$

If U > 1, nothing is feasible

If $r_i = 0$ and $p_i = d_i$, then write $T_i = (p_i, e_i)$

Types of Scheduling

Static

Static Scheduling:

- task's priority is assigned before execution and does not change
- If a task misses its deadline, you mess up all the deadlines after it like an airport at Christmas

FIFO

First In First Out (FIFO):

- Could cause problems for tasks whose execution time is significantly shorter than the rest when there are deadlines
    - E.g. $T_1 = (100, 3)$; $T_2 = (2, 1)$
- A.K.A. **First Come, First Served (FCFS)**

Cyclic Executive: frame-based scheduling

- When you allocate an amount of time where a task can execute
- Can have multiple executions of the same task
- Tasks might not even fill the full frame

Schedule: the order in which tasks will be executed
Hyperperiod [H]: the entire length of a cycle, least common multiple

Harmonic: every task period evenly divides every longer period

Pre-empting: splitting a task up into multiple mini tasks. Also, if a task misses its deadline, halt the task at the deadline

Frame Size [f]:

- The best way for computers to segment the schedule in a way that it verify that the appropriate tasks have been executed
- Process: try each see which is the largest frame size that follows all the below constraints from 1 to e_{max}.
- Constraints:
  1. $f \geq \max_{1 \leq i \leq n}(e_i)$
  2. H % f = 0
  3. $2f - \gcd(p_i, f) \leq d_i$

Least Compute Time (LCT): tasks with smallest execution times executed first

- Think *greedy*
- Works poorly; worse than RR

Rate Monotonic (RM): shorter period, higher priority

- Think: tasks requiring frequent attention should have higher priority
- If harmonic, feasible as long as U ≤ 1
- If non-harmonic, guaranteed feasible if $U \leq n\left(2^{\frac{1}{n}} - 1\right)$
  - If the equation fails, it still might be, so draw the whole thing to be safe.

Dynamic

Dynamic: each of the tasks' priorities can change. *Think*: while for static priorities it is constantly re-evaluating <u>which</u> task has the highest priority, dynamic scheduling also re-evaluates <u>the actual</u> priorities, themselves.

The only two optimal dynamic priorities are:

- **Earliest Deadline First (EDF)**:
  - more flexible, better U
  - If deadlines < periods, still optimal, but determining feasibility is NP-hard
  - Always feasible if U ≤ 1
- **Least Slack Theorem (LST)**: not as popular as EDF

Multiprocessor

Once you have multiple processors, neither EDF nor RM are guaranteed to work.

Look into first-fit algorithms

Task Interactions

Suspended: active choice, of access prevention until algorithm allows it to
Blocked: as a result of waiting for a resource to be free

How to do the timing diagrams with locks:

- S_1 = lock(S_1)
- S_1^ = unlock(S_1)

One-shot Tasks: non-periodic tasks

Critical Section: when a task tries to acquire an already locked by another task resource

Priority Inversion: a method of avoiding deadlock by telling high priority tasks to share their resources with the lower priority tasks even when it's not their turn
- Allocate time, where T_1 has access to shared resource, so the time not allocated can be pre-empted
- Connect the pre-empted by T_1 when T_1 wants to access the resource
- Protect the resource with a semaphore
- You can make it so that tasks can use the resource even after they release the semaphore, but you risk overwriting in that time

Priority Inheritance Protocol (PIP):
- Temporarily raise the priority of a task only if and when it actually blocks a higher priority task; on leaving the critical section, the task priority reverts to its original value
- Issues:
  - If only one shared resource, there's only one possible schedule
  - If more than one resource blocking:
    - Blocking time may be excessively long
    - Deadlock may occur
  - If accessing multiple resources, you can only use them in the same order

Priority Ceiling Protocol (PCP): tasks entering a critical section can only access the blocked resource if it has a priority higher than the priority ceiling
- **Priority Ceiling (PC)**: maximum priority of all tasks ever going to access a resource
- Only need to check PC when entering a critical section
- If any task needs priority higher than the priority ceiling, it's suspended
- Main advantages:
  - No locked resources, so free access
  - "The state of the art when resolving resource-contention issues"
  - "Deadlock free for an arbitrary number of tasks with an arbitrary number of resources acted upon in an arbitrary way."

Sporadic Server
Execution Budget [e_s]: periodic tasks aren't flexible…
Execution time [e_i]: …sporadic tasks are

Deadline [d_i]: absolute deadline
Release Time [r_i]:

Set of Sporadic Tasks [θ]:
Sporadic Task [S_i]:
- Non-periodic task

- (r_i, e_i, d_i)
- Typically interrupt-driven

Rules [ρ]: set of rules regulating a sporadic server

Sporadic Server [Φ_s]: (p_s, e_s, θ, ρ)
Periodic Task: (p_s, e_s) a type of sporadic server

- no expectation of when it finishes, only that a new one is queued every period

Assume:

- Φ_s scheduled with T_i according to RM

We don't use K_d because it looks at the derivative regardless of the size of the error function. If your error is a sine function with a small amplitude, K_d will only take the derivative into account and it will overcompensate.

Open loop response: plant with no control

Ziegler-Nichols Tuning Rule: a PID tuning rule
Look at the *open loop response*. It could have a longer rise time / overshoot than preferred.

1. Tangent to curve on upslope

High sample rate → lots of high frequency noise

Effective Utilization [δ]:

$U = U_{periodic} + \delta U_{sporadic}$

Error bound [ε]:

Slack [ω]:
Acceptance Test: check of stuff

$$\omega(S_k, t) = \left\lfloor \frac{d_k - t}{p_s} \right\rfloor e_s - e_k - \sum_{S_i \in \theta : d_i < d_k} e_i - \xi_i$$

1. If $\omega(S_k, t) < 0$, reject task
2. If $\omega(S_k, t) \geq 0$, need to check if already accepted sporadic tasks are adversely affected, i.e. $\omega(S_j, t) - e_k \geq 0$ holds for all $S_j \in \theta$ with $d_j \geq d_k$.

The set θ is maintained dynamically.

Clocks

Computer Clock [C]:

Attributes:

- Correctness
- Bounded Drift
- Monotonicity

- Chronoscopicity

Drift [ρ]: rate of change of the clock value away from a perfect clock (each second)
There's usually a reason why a clock drifts

$$\left|\frac{\mathrm{d}C(t)}{\mathrm{d}t}-1\right|\leq\rho$$

(EPS):

Monotonicity: Clock will always have a consistent spacing and will only move in one order (forward / backwards)

SSL certs will fail signature if your clock is wrong as to ensure this

Chronoscopicity [γ]: changing drift
second derivative of stuff $\left|\frac{\mathrm{d}^2C(t)}{\mathrm{d}t^2}\right|\leq\gamma$

PID Control

Plant [G(s)]: a transfer function, e.g. $\frac{1}{s^2+10s+20}$

Remember this from 3DX4? Most of the stuff is still there, so refer to that. More here.
Each of the K's represent a different error or gain

4 types of controllers [P(s)]:

- **Proportional Controller (P),(PC)**: $\frac{K_p}{s^2+10s+\left(20+K_p\right)}$
- **Proportional Integral (PI)**: $\frac{K_ps+K_i}{s^3+10s^2+\left(20+K_ps+K_i\right)}$
- **Proportional Derivative (PD)**: $\frac{K_ds+K_p}{s^2+\left(10+K_d\right)s+\left(20+K_p\right)}$
- **Proportional Integral Derivative (PID)**: $\frac{K_ps+K_i+s^2K_d}{s\left(s^2+as+b\right)+K_ps+K_i+s^2K_d}$

  $$\frac{K_ds^2+K_ps+K_i}{s^3+\left(10+K_d\right)s^2+\left(20+K_p\right)s+K_i}$$

$$u(t)=K_pe(t)+K_d\dot{e}(t)+K_i\int_0^t e(v)\mathrm{d}v$$

Designing a PID Controller

1. Obtain an open-loop response and determine what needs to be improved
2. Add K_p to improve the rise time
3. Add K_d to improve the overshoot
4. Add K_i to eliminate the steady-state error
5. Adjust each of K_p, K_d, and K_i until you obtain a desired overall response. You can always refer to the table below to find out which controller controls what characteristics.

Increasing this	RISE TIME	OVERSHOOT	SETTLING TIME	S-S ERROR
K_p	Decrease	Increase	Small Change	Decrease
K_d	Small Change	Decrease	Decrease	No Change
K_i	Decrease	Increase	Increase	Eliminate

Ziegler-Nichols Tuning Rule:

- a plant with neither integrators nor dominant complex-conjugate pairs
- Look at the *open loop response*. It could have a longer rise time / overshoot than preferred.
- Tangent to curve on upslope
- For PID controllers

Noise frequency [ω_n]:
Noise amplitude [a_n]:
Open loop: plant with no control
Closed loop: $\frac{C(s)G(s)}{1+C(s)G(s)}$

Jitter

Jitter: a delay