

SFWR ENG 4AA4

Kemal Ahmed
Fall 2015
Dr. Down

Note: information from the pre-requisite, [SFWR ENG 3DX4](#) will not be included in this summary (although corrections will be).

Contents

Real-Time Systems	1
Classifications.....	1
Task optimization	2
Types of Scheduling	2
Static	2
FIFO	2
Dynamic	3
Multiprocessor	3
Task Interactions	4
Sporadic Server	5
Clocks	6
Christian's Algorithm.....	7
Berkeley	7
PID Control	7
Designing a PID Controller	8
Jitter	8

Real-Time Systems

Classifications

What happens upon failure to meet deadlines:

- **Soft:** performance is degraded but not destroyed
- **Firm:** a few times will simply degrade performance, but after may lead to system failure
- **Hard:** complete and catastrophic system failure
 - **Safety Critical:** may cause injury / death (a type of hard)

Forward difference method: derivatives using $f'(x) = \frac{f(x+h) - f(x)}{h}$

Backwards Difference method: derivatives using $f'(x) = \frac{f(x) - f(x-h)}{h}$

Controller [C(s)]:

Input [E(s)]:

Output [U(s)]:

$$U(s) = C(s)E(s)$$

Task optimization

Task [T]: $T_i = (p_i, r_i, e_i, d_i)$

Period [p]: time between tasks are repeatedly released

Release time [r]: time it takes to release task

Execution time [e]: slowest time task could take to be completed (but assume the tasks will take this long no matter what)

Deadline [d]: when task needs to be completed

Number of tasks [n]:

Processor Utilization [U]: used as a priority level $U = \sum_{i=1}^n \frac{e_i}{p_i}$

If $U > 1$, nothing is feasible

If $r_i = 0$ and $p_i = d_i$, then write $T_i = (p_i, e_i)$

Types of Scheduling

Static

Static Scheduling:

- task's priority is assigned before execution and does not change
- If a task misses its deadline, you mess up all the deadlines after it like an airport at Christmas
- A.K.A. **Fixed priority**

FIFO

First In First Out (FIFO):

- Could cause problems for tasks whose execution time is significantly shorter than the rest when there are deadlines
 - E.g. $T_1 = (100, 3)$; $T_2 = (2, 1)$
- A.K.A. **First Come, First Served (FCFS)**

Cyclic Executive: frame-based scheduling

- When you allocate an amount of time where a task can execute
- Can have multiple executions of the same task
- Tasks might not even fill the full frame

Schedule: the order in which tasks will be executed

Hyperperiod [H]: the entire length of a cycle, least common multiple

Harmonic: every task period evenly divides every longer period

Pre-empting: splitting a task up into multiple mini tasks. Also, if a task misses its deadline, halt the task at the deadline

Frame Size [f]:

- The best way for computers to segment the schedule in a way that it verify that the appropriate tasks have been executed
- Process: try each see which is the largest frame size that follows all the below constraints from 1 to e_{\max} .
- Constraints:
 1. $f \geq \max_{1 \leq i \leq n} (e_i)$
 2. $H \% f = 0$
 3. $2f - \gcd(p_i, f) \leq d_i$

Least Compute Time (LCT): tasks with smallest execution times executed first

- Think *greedy*
- Works poorly; worse than RR

Rate Monotonic (RM): shorter period, higher priority

- Think: tasks requiring frequent attention should have higher priority
- If harmonic, feasible as long as $U \leq 1$
- If non-harmonic, guaranteed feasible if $U \leq n(2^{\frac{1}{n}} - 1)$
 - If the equation fails, it still might be, so draw the whole thing to be safe.

Dynamic

Dynamic: each of the tasks' priorities can change. *Think:* while for static priorities it is constantly re-evaluating which task has the highest priority, dynamic scheduling also re-evaluates the actual priorities, themselves.

The only two optimal dynamic priorities are:

- **Earliest Deadline First (EDF):**
 - more flexible, better U
 - If deadlines < periods, still optimal, but determining feasibility is NP-hard
 - Always feasible if $U \leq 1$
- **Least Slack Theorem (LST):** not as popular as EDF

Multiprocessor

Once you have multiple processors, neither EDF nor RM are guaranteed to work.

Look into first-fit algorithms

Task Interactions

Suspended: active choice, of access prevention until algorithm allows it to

Blocked: as a result of waiting for a resource to be free

How to do the timing diagrams with locks:

- $S_1 = \text{lock}(S_1)$
- $S_1^{\wedge} = \text{unlock}(S_1)$

One-shot Tasks: non-periodic tasks

Critical Section: when a task tries to acquire a shared resource already locked by another task

Priority Inversion: a method of avoiding deadlock by telling high priority tasks to share their resources with the lower priority tasks even when it's not their turn

- Allocate time, where T_1 has access to shared resource, so the time not allocated can be pre-empted
- Connect the pre-empted by T_1 when T_1 wants to access the resource
- Protect the resource with a semaphore
- You can make it so that tasks can use the resource even after they release the semaphore, but you risk overwriting in that time

Priority Inheritance Protocol (PIP):

- Temporarily raise the priority of a task only if and when it actually blocks a higher priority task; on leaving the critical section, the task priority reverts to its original value
- Issues:
 - If only one shared resource, there's only one possible schedule
 - If more than one resource blocking:
 - Blocking time may be excessively long
 - Deadlock may occur
 - If accessing multiple resources, you can only use them in the same order

Priority Ceiling Protocol (PCP): tasks entering a critical section can only access the blocked resource if it has a priority higher than the priority ceiling

- **Priority Ceiling (PC):** maximum priority of all tasks ever going to access a resource
- Only need to check PC when entering a critical section
- If any task needs priority higher than the priority ceiling of ALL of the semaphores currently locked, it's suspended
- Main advantages:
 - No locked resources, so free access
 - "The state of the art when resolving resource-contention issues"
 - "Deadlock free for an arbitrary number of tasks with an arbitrary number of resources acted upon in an arbitrary way."
 - **Deadlock:** think if you and I are at a table with one fork and one knife and you need both to eat, but you take the fork and I take the knife.

Sporadic Server

Execution Budget [e_s]: periodic tasks aren't flexible...

Execution time [e_i]: ...sporadic tasks are

Deadline [d_i]: absolute deadline

Release Time [r_i]:

Set of Sporadic Tasks [θ]:

Sporadic Task [S_i]:

- Non-periodic task
- (r_i, e_i, d_i)
- Typically interrupt-driven

Rules [ρ]: set of rules regulating a sporadic server

Sporadic Server [Φ_s]: (p_s, e_s, θ, ρ)

Periodic Task: (p_s, e_s) a type of sporadic server

- no expectation of when it finishes, only that a new one is queued every period

Assume:

- Φ_s scheduled with T_i according to RM

We don't use K_d because it looks at the derivative regardless of the size of the error function. If your error is a sine function with a small amplitude, K_d will only take the derivative into account and it will overcompensate.

Open loop response: plant with no control

Ziegler-Nichols Tuning Rule: a PID tuning rule

Look at the *open loop response*. It could have a longer rise time / overshoot than preferred.

1. Tangent to curve on upslope

High sample rate \rightarrow lots of high frequency noise

Effective Utilization [δ]:

$$U = U_{\text{periodic}} + \delta U_{\text{sporadic}}$$

Error bound [ϵ]:

Slack [ω]:

Acceptance Test: check of stuff

$$\omega(S_k, t) = \left\lfloor \frac{d_k - t}{p_s} \right\rfloor e_s - e_k - \sum_{S_i \in \theta: d_i < d_k} e_i - \xi_i$$

1. If $\omega(S_k, t) < 0$, reject task

2. If $\omega(S_k, t) \geq 0$, need to check if already accepted sporadic tasks are adversely affected, i.e. $\omega(S_j, t) - e_k \geq 0$ holds for all $S_j \in \theta$ with $d_j \geq d_k$.

The set θ is maintained dynamically.

Clocks

Computer Clock [C]:

Standard Clock [C_s]: perfect clock; has real time

Attributes:

- Correctness
- Bounded Drift
- Monotonicity
- Chronoscopia

(EPS): a bounded/maximum difference between the clock time and the real time

$$|C(t) - C_s(t)| \leq EPS$$

Reset time [r]: the real time you set the clock to when you reset it

Drift [E]: rate of change of the clock value away from a perfect clock (each second)

- There's usually a reason why a clock drifts

Drift Bound [ρ]: maximum drift

$$\left| \frac{dC(t)}{dt} - 1 \right| \leq \rho$$

Reset Error [ε]: error between actual time and time clock was set to at reset

Total Error [E]: $E(t) = \epsilon + \text{drift_since_reset}$

$$\text{drift_since_reset} \leq \rho(C(t) - r)$$

$$E(t) = \rho(C(t) - r) + \epsilon \leq EPS$$

$$C(t) - r \leq (EPS - \epsilon)/\rho$$

Real time will be within this interval – $[C(t) - E(t), C(t) + E(t)]$

Monotonicity: Clock will always have a consistent spacing and will only move in one order (forward / backwards)

SSL certs will fail signature if your clock is wrong as to ensure this

Chronoscopia [γ]: maximum changing drift

$$\text{second derivative of stuff } \left| \frac{d^2C(t)}{dt^2} \right| \leq \gamma$$

Christian's Algorithm

Minimum Latency [T_{\min}]:

Request Send Time [T_0]:

Request Receive Time [T_1]:

Server Time [T_{server}]: time returned by the server


$$T_{\text{new}} = T_{\text{server}} + \frac{T_1 - T_0}{2}$$

$$\text{Accuracy is } \pm \frac{T_1 - T_0}{2} - T_{\min}$$

Berkeley

Not often used, but useful for learning

1. Elect 1 node to be the **master**, the one that runs the algorithm



2. Finds the average of the nodes. However, that's probably going to find a value that isn't near any of them.
3. Eliminate the outliers:
 - a. Standard deviation: the more outliers, the harder to remove them
 - b. Median
 - c. **Maximum deviation**: maximum clock drift × time since last synchronization; sometimes it's good to use physical limitations as the minimum check to ensure accuracy

PID Control

Plant [$G(s)$]: a transfer function, e.g. $\frac{1}{s^2 + 10s + 20}$

Remember this from 3DX4? Most of the stuff is still there, so refer to that. [More here](#).

Each of the K's represent a different error or gain

4 types of controllers [$P(s)$]:

- **Proportional Controller (P),(PC)**: $\frac{K_p}{s^2 + 10s + (20 + K_p)}$
- **Proportional Integral (PI)**: $\frac{K_p s + K_i}{s^3 + 10s^2 + (20 + K_p s + K_i)}$
- **Proportional Derivative (PD)**: $\frac{K_d s + K_p}{s^2 + (10 + K_d)s + (20 + K_p)}$

- **Proportional Integral Derivative (PID):** $\frac{K_p s + K_i + s^2 K_d}{s(s^2 + as + b) + K_p s + K_i + s^2 K_d}$

$$\frac{K_d s^2 + K_p s + K_i}{s^3 + (10 + K_d)s^2 + (20 + K_p)s + K_i}$$

$$u(t) = K_p e(t) + K_d \dot{e}(t) + K_i \int_0^t e(v) dv$$

Designing a PID Controller

1. Obtain an open-loop response and determine what needs to be improved
2. Add K_p to improve the rise time
3. Add K_d to improve the overshoot
4. Add K_i to eliminate the steady-state error
5. Adjust each of K_p , K_d , and K_i until you obtain a desired overall response. You can always refer to the table below to find out which controller controls what characteristics.

Increasing this	RISE TIME	OVERSHOOT	SETTLING TIME	S-S ERROR
K_p	Decrease	Increase	Small Change	Decrease
K_d	Small Change	Decrease	Decrease	No Change
K_i	Decrease	Increase	Increase	Eliminate

Ziegler-Nichols Tuning Rule:

- a plant with neither integrators nor dominant complex-conjugate pairs
- Look at the *open loop response*. It could have a longer rise time / overshoot than preferred.
- Tangent to curve on upslope
- For PID controllers

Noise frequency $[\omega_n]$:

Noise amplitude $[a_n]$:

Open loop: plant with no control

Closed loop: $\frac{C(s)G(s)}{1 + C(s)G(s)}$

Jitter

Jitter [J]: a delay

Relative Jitter: difference in response time between current and previous response times

$$\max_k |R_{i,k+1} - R_{i,k}|$$

Absolute Jitter: difference between largest response time and smallest response time

$$\max_k R_{i,k} - \min_k R_{i,k}$$

Absolute jitter \geq relative jitter